


number of background workers on the common_site_config.json. Gunicorn workers are different (they handle the requests).

In development mode, usually the # of worker is set to 1. **Scheduler** (like scheduler events) **is only responsible for scheduling the work/job** (when the time comes, it will tell the worker to do the job). **The one that is responsible to do the actual work/job is the background workers.**



In Frappe, there are 3 types of task/queue: short, default, long

Now, we can define which types of task that a worker is supposed to take. For example, if we have 2 workers, we can set it in such a way so that **worker 1** will only take **short** tasks, while **worker 2** will only take **default** & **long** tasks.

To handle these queueing, frappe uses **redis queue**. We can see the worker on the **RQ Worker** virtual doctype that comes from Redis directly, it tracks the background worker that we have:


| ID | Status | Queue Type(s) | Failed Job Count | Utilization % | Last Updated On |
|-------|--------|----------------------|------------------|---------------|-----------------|
| 89638 | idle | short, default, long | 15 | 15% | now |

From here, we can see the worker we have, what status it's currently in, the queue type it handles, etc. Utilization % means that if it was running for 1 hour, how many of that 1 hour is it actually working.

Similarly, frappe also tracks the background jobs on **RQ Job** doctype:

| ID | Job Name | Queue | Status | Queue | Last Updated On |
|----|--|---------|----------|---------|-----------------|
| | frappe.utils.change_log.check_for_update | long | failed | long | 1h |
| | frappe.twofactor.delete_all_barcode_for_users | default | finished | default | 6m |
| | frappe.website.doctype.web_page.web_page.check | default | finished | default | 6m |
| | frappe.oauth.delete_oauth2_data | default | finished | default | 6m |
| | frappe.automation.doctype.reminder.reminder.send | default | finished | default | 3m |
| | frappe.monitor.flush | default | finished | default | 3m |
| | frappe.utils.global_search.sync_global_search | default | finished | default | 3m |
| | frappe.email.doctype.email_account.email_account | default | finished | default | 3m |
| | frappe.email.queue.flush | default | finished | default | 3m |
| | frappe.email.doctype.email_account.email_account | default | finished | default | 1m |

Default timeout for **default** is 5 minutes, while for **long** it's 25 minutes.



There is also this thing called **batching**, where we can break down a big job into multiple smaller jobs so that we can work on them in parallel.

How to queue our jobs:

We can use the method `frappe.enqueue()` and provide the dotted path to the method.

```

api.py 1, M x
cabs_app > api.py > ...
9      )
10     return sum(amounts)
11
12     def takes_a_long_time():
13         import time
14         time.sleep(15)
    You, 25 seconds ago • Uncommitt

PROBLEMS 1 TERMINAL OUTPUT GITLENS
> v TERMINAL Python - be
  o → bench-0 bench --site irfancabs.local console
    Apps in this namespace:
    frappe, cabs_app

    In [1]: frappe.enqueue("cabs_app.api.takes_a_long_time")
  
```

If we run the command above, we can see the job being queued and worked on the `RQ Job` virtual doctype. We can also tag the queue by adding the argument `queue="short"` (or other type).

To add workers, we can just add new entries for the worker process on **Procfile** and give them different names:

```

py 1, M Procfile x
> mdhussain > Frappe > bench-0 > Procfile
Schedule: bench schedule

worker: bench worker 1>> logs/worker.log 2>> logs/worker.error.log
worker-1: bench worker 1>> logs/worker.log 2>> logs/worker.error.log
worker-2: bench worker 1>> logs/worker.log 2>> logs/worker.error.log
  
```